

# x86 Memory management

*What was it you were talking about?*

© Government of Canada

This document is the property of the Government of Canada. It shall not be altered, distributed beyond its intended audience, produced, reproduced or published, in whole or in any substantial part thereof, without the express permission of CSE.



Communications  
Security Establishment

Centre de la sécurité  
des télécommunications

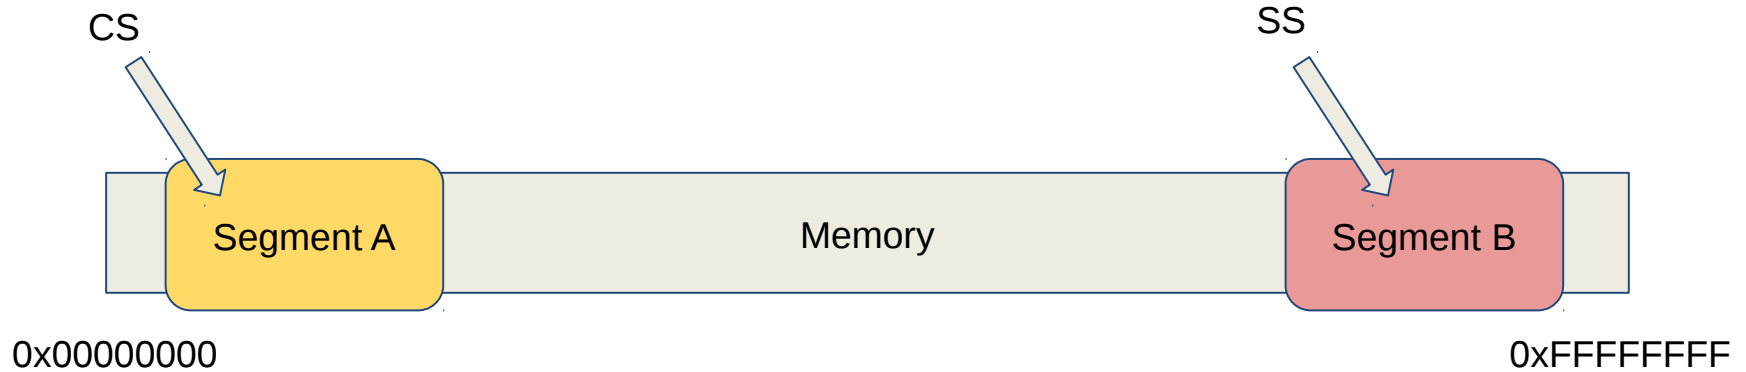
PAGE 1

Canada 

So, exactly what is the difference  
between segmentation and virtual  
memory?




# What is segmentation?



A segment defines a **zone of memory**. A segment register then points to that zone allowing access to it.

# Segmentation setup x86



The GDTR or LDTR allow to find the Global descriptor table or Local descriptor table on the memory.

GDTR / IDTR:

- Base: 0x1000
- Limit: 0x18

# On Linux

## traps.c

```

1015 ▼ #ifdef CONFIG_X86_32
1016 »     set_system_intr_gate(IA32_SYSCALL_VECTOR, entry_INT80_32);
1017 »     set_bit(IA32_SYSCALL_VECTOR, used_vectors);
1018 #endif
1019
1020 ▼ »     /*
1021 »     * Set the IDT descriptor to a fixed read-only location, so that the
1022 »     * "sidt" instruction will not leak the location of the kernel, and
1023 »     * to defend the IDT against arbit
1024 »     * It will be reloaded in cpu_init
1025 »     _set_fixmap(FIX_RO_IDT, __pa_symbol
1026 »     idt_descr.address = fix_to_virt(FI
1027
1028 ▼ »     /*
1029 »     * Should be a barrier for any ext
1030 »     */
1031 »     cpu_init();

```

```

1576 void cpu_init(void)
1577 ▼ {
1578 »     int cpu = smp_processor_id();
1579 »     struct task_struct *curr = current;
1580 »     struct tss_struct *t = &per_cpu(cpu_tss, cpu);
1581 »     struct thread_struct *thread = &curr->thread;
1582
1583 »     wait_for_master_cpu(cpu);
1584
1585 ▼ »     /*
1586 »     * Initialize the CR4 shadow before doing anything that could
1587 »     * try to read it.
1588 »     */
1589 »     cr4_init_shadow();
1590
1591 »     show_ucode_info_early();
1592
1593 »     pr_info("Initializing CPU%d\n", cpu);
1594
1595 »     if (cpu_feature_enabled(X86_FEATURE_VME) ||
1596 »         boot_cpu_has(X86_FEATURE_TSC) ||
1597 »         boot_cpu_has(X86_FEATURE_DE))
1598 »         cr4_clear_bits(X86_CR4_VME|X86_CR4_PVI|X86_CR4_TSD|X86_CR4_DE);
1599
1600 »     load_current_idt();
1601 »     switch_to_new_gdt(cpu);
1602

```

## common.c



# On Linux

## common.c

```

500 void switch_to_new_gdt(int cpu)
501 {
502     /* Load the original GDT */
503     load_direct_gdt(cpu);
504     /* Reload the per-cpu base */
505     load_percpu_segment(cpu);
506 }
507
475 void load_direct_gdt(int cpu)
476 {
477     struct desc_ptr gdt_descr;
478
479     gdt_descr.address = (long)get_cpu_gdt_rw(cpu);
480     gdt_descr.size = GDT_SIZE - 1;
481     load_gdt(&gdt_descr);
482 }

```

## desc.h

```

43 struct gdt_page {
44     struct desc_struct gdt[GDT_ENTRIES];
45 } __attribute__((aligned(PAGE_SIZE)));
46
47 DECLARE_PER_CPU_PAGE_ALIGNED(struct gdt_page, gdt_page);
48
49 /* Provide the original GDT */
50 static inline struct desc_struct *get_cpu_gdt_rw(unsigned int cpu)
51 {
52     return per_cpu(gdt_page, cpu).gdt;
53 }
54
125 #define load_gdt(dtr) >> >> >> native_load_gdt(dtr)
236 static inline void native_load_gdt(const
237 {
238     asm volatile("lgdt %0"::"m" (*dtr)
239 }

```

# On Linux

Bolded bits shows the DPL (descriptor privilege level).

0 = Kernel mode, 3 = User mode.

Definite proof that Linux only uses 2 of the 4 rings available on x86!

## common.c

```

101 DEFINE_PER_CPU_PAGE_ALIGNED(struct gdt_page, gdt_page) = { .gdt = {
102 #ifdef CONFIG_X86_64
103     /*
104      * We need valid kernel segments for data and code in long mode too
105      * IRET will check the segment types kkeil 2000/10/28
106      * Also sysret mandates a special GDT layout
107      *
108      * TLS descriptors are currently at a different place compared to i386.
109      * Hopefully nobody expects them at a fixed place (Wine?)
110      */
111     [GDT_ENTRY_KERNEL32_CS] = GDT_ENTRY_INIT(0xc09b, 0, 0xffff),
112     [GDT_ENTRY_KERNEL_CS] = GDT_ENTRY_INIT(0xa09b, 0, 0xffff),
113     [GDT_ENTRY_KERNEL_DS] = GDT_ENTRY_INIT(0xc093, 0, 0xffff),
114     [GDT_ENTRY_DEFAULT_USER32_CS] = GDT_ENTRY_INIT(0xc0fb, 0, 0xffff),
115     [GDT_ENTRY_DEFAULT_USER_DS] = GDT_ENTRY_INIT(0xc0f3, 0, 0xffff),
116     [GDT_ENTRY_DEFAULT_USER_CS] = GDT_ENTRY_INIT(0xa0fb, 0, 0xffff),
117 #else
118     [GDT_ENTRY_KERNEL_CS] = GDT_ENTRY_INIT(0xc09a, 0, 0xffff),
119     [GDT_ENTRY_KERNEL_DS] = GDT_ENTRY_INIT(0xc092, 0, 0xffff),
120     [GDT_ENTRY_DEFAULT_USER_CS] = GDT_ENTRY_INIT(0xc0fa, 0, 0xffff),
121     [GDT_ENTRY_DEFAULT_USER_DS] = GDT_ENTRY_INIT(0xc0f2, 0, 0xffff),

```

Looking at the flags for the 32 bits version, you should see the following binary pattern for Kernel mode code:

0b 1100 0000 **1001** 1010

Looking at the flags for the 32 bits version, you should see the following binary pattern for User mode code:

0b 1100 0000 **1111** 1010

```

14 /*
15  * FIXME: Accessing the desc_struct through its fields is more elegant,
16  * and should be the one valid thing to do. However, a lot of open code
17  * still touches the a and b accessors, and doing this allow us to do it
18  * incrementally. We keep the signature as a struct, rather than a union,
19  * so we can get rid of it transparently in the future -- glommer
20  */
21 /* 8 byte segment descriptor */
22 struct desc_struct {
23     union {
24         struct {
25             unsigned int a;
26             unsigned int b;
27         };
28         struct {
29             u16 limit0;
30             u16 base0;
31             unsigned base1: 8, type: 4, s: 1, dpl: 2, p: 1;
32             unsigned limit: 4, avl: 1, l: 1, d: 1, g: 1, base2: 8;
33         };
34     };
35 } __attribute__((packed));
36
37 #define GDT_ENTRY_INIT(flags, base, limit) { { { \
38     .a = (((limit) & 0xffff) | (((base) & 0xffff) << 16)), \
39     .b = (((base) & 0xff0000) >> 16) | (((flags) & 0xf0ff) << 8) | \
40     ((limit) & 0xf0000) | ((base) & 0xff000000), \
41     } } }
42

```

desc\_defs.h

Alright, let's write an OS...  
Just kidding, let's talk about virtual  
memory...



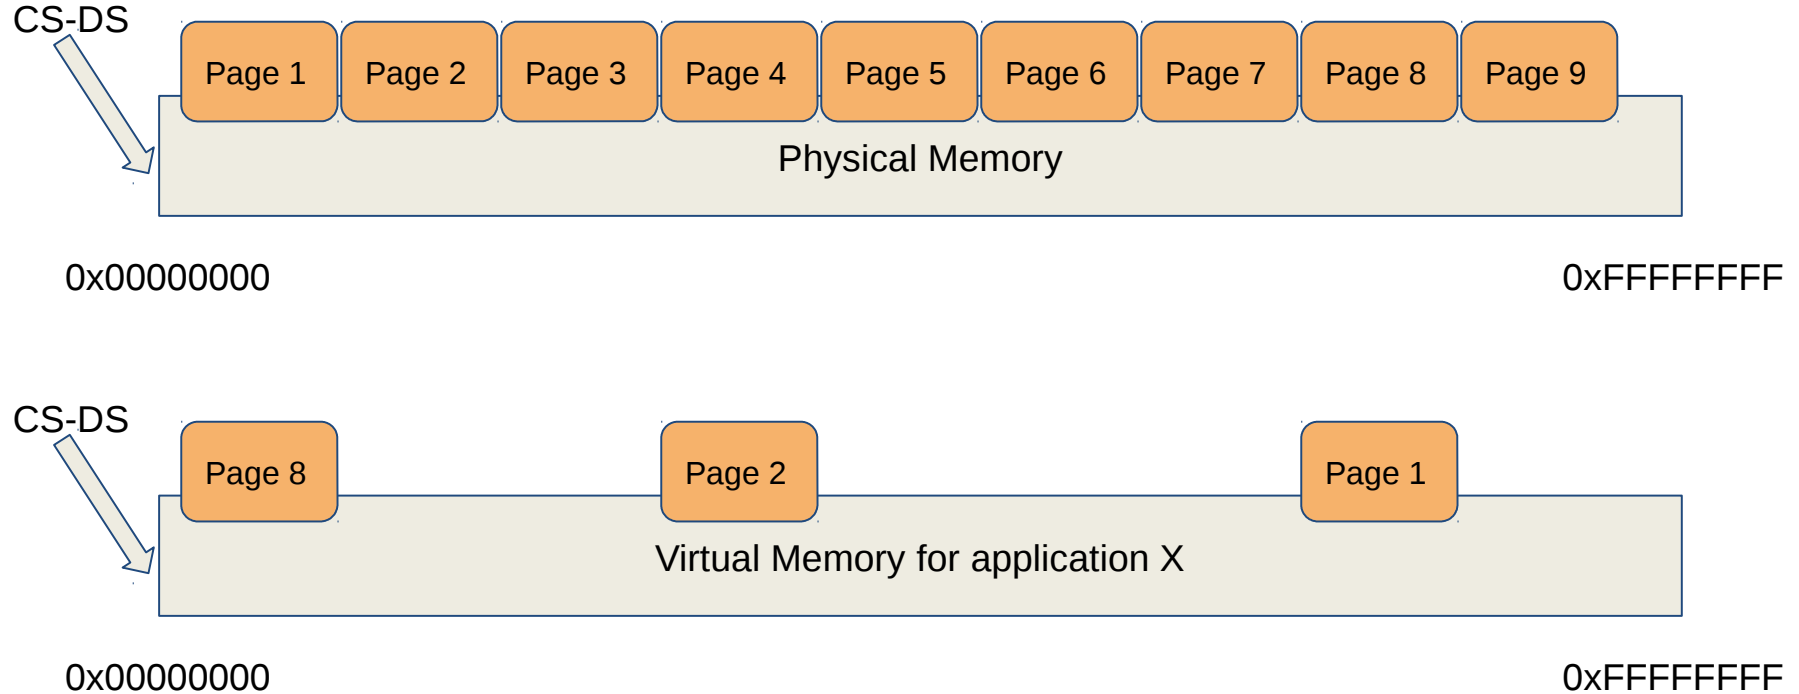


On x86, granularity in memory management is made possible by the use of virtual memory.

Be careful, Intel calls that “paging”. Not to be confused with paging/swapping.

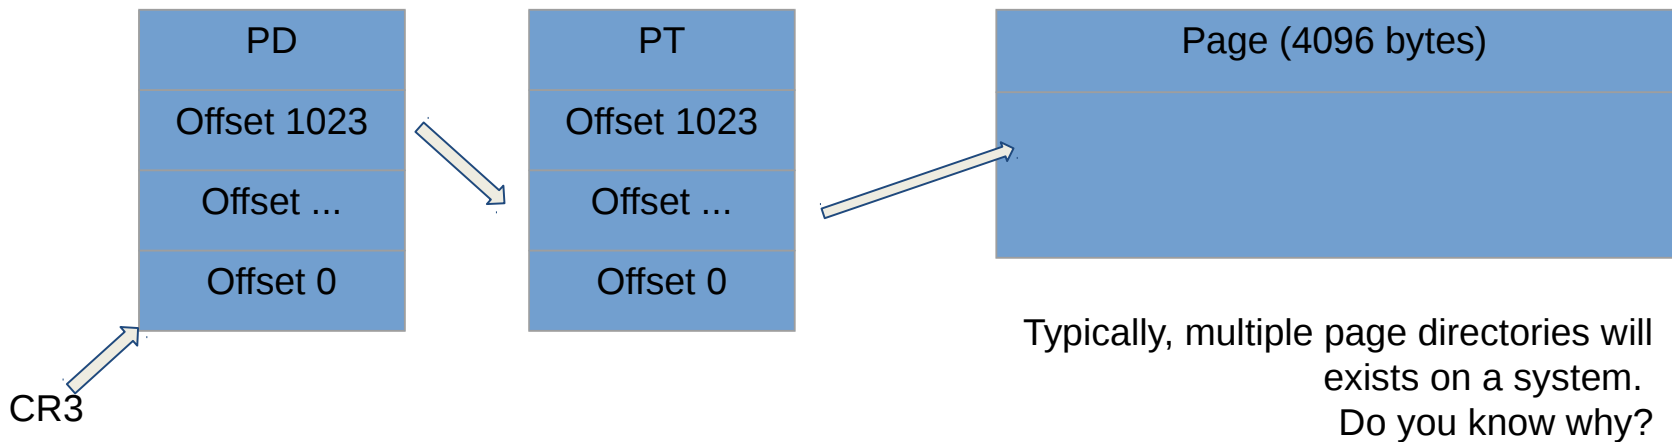


# Virtual memory



# Virtual Address

31 - 22	21 - 12	11 - 0
Page Directory Offset	Page Table Offset	Page Offset



# On Linux

## head\_32.s

```
128  #endif
129
130  >      /* Create early pagetables. */
131  >      call mk_early_pgtbl_32
132
290  /*
291  * Enable paging
292  */
293  >      movl $pa(initial_page_table), %eax
294  >      movl %eax,%cr3> >      /* set the page table pointer.. */
295  >      movl $CR0_STATE,%eax
296  >      movl %eax,%cr0> >      /* ..and set paging (PG) bit */
297  >      ljmp $__BOOT_CS,$1f>      /* Clear prefetch and normalize %eip */
298  1:
299  >      /* Shift the stack pointer to a virtual address */
300  >      addl $__PAGE_OFFSET, %esp
```

Before line 296, the system is running using physical addresses. Once x86 paging is activated, the system is running on virtual address.

The jump and esp modifications are there to make sure pointers shows virtual addresses.