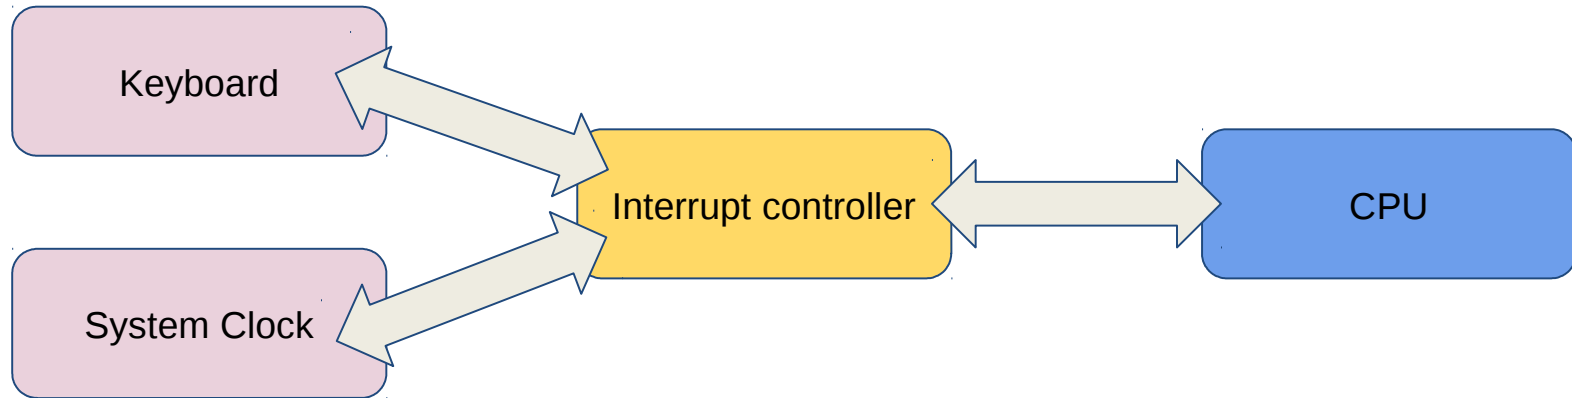


Interrupt Management


We are interrupting this presentation so we can have a quick chat about x86 interrupt management.



High level view of the system

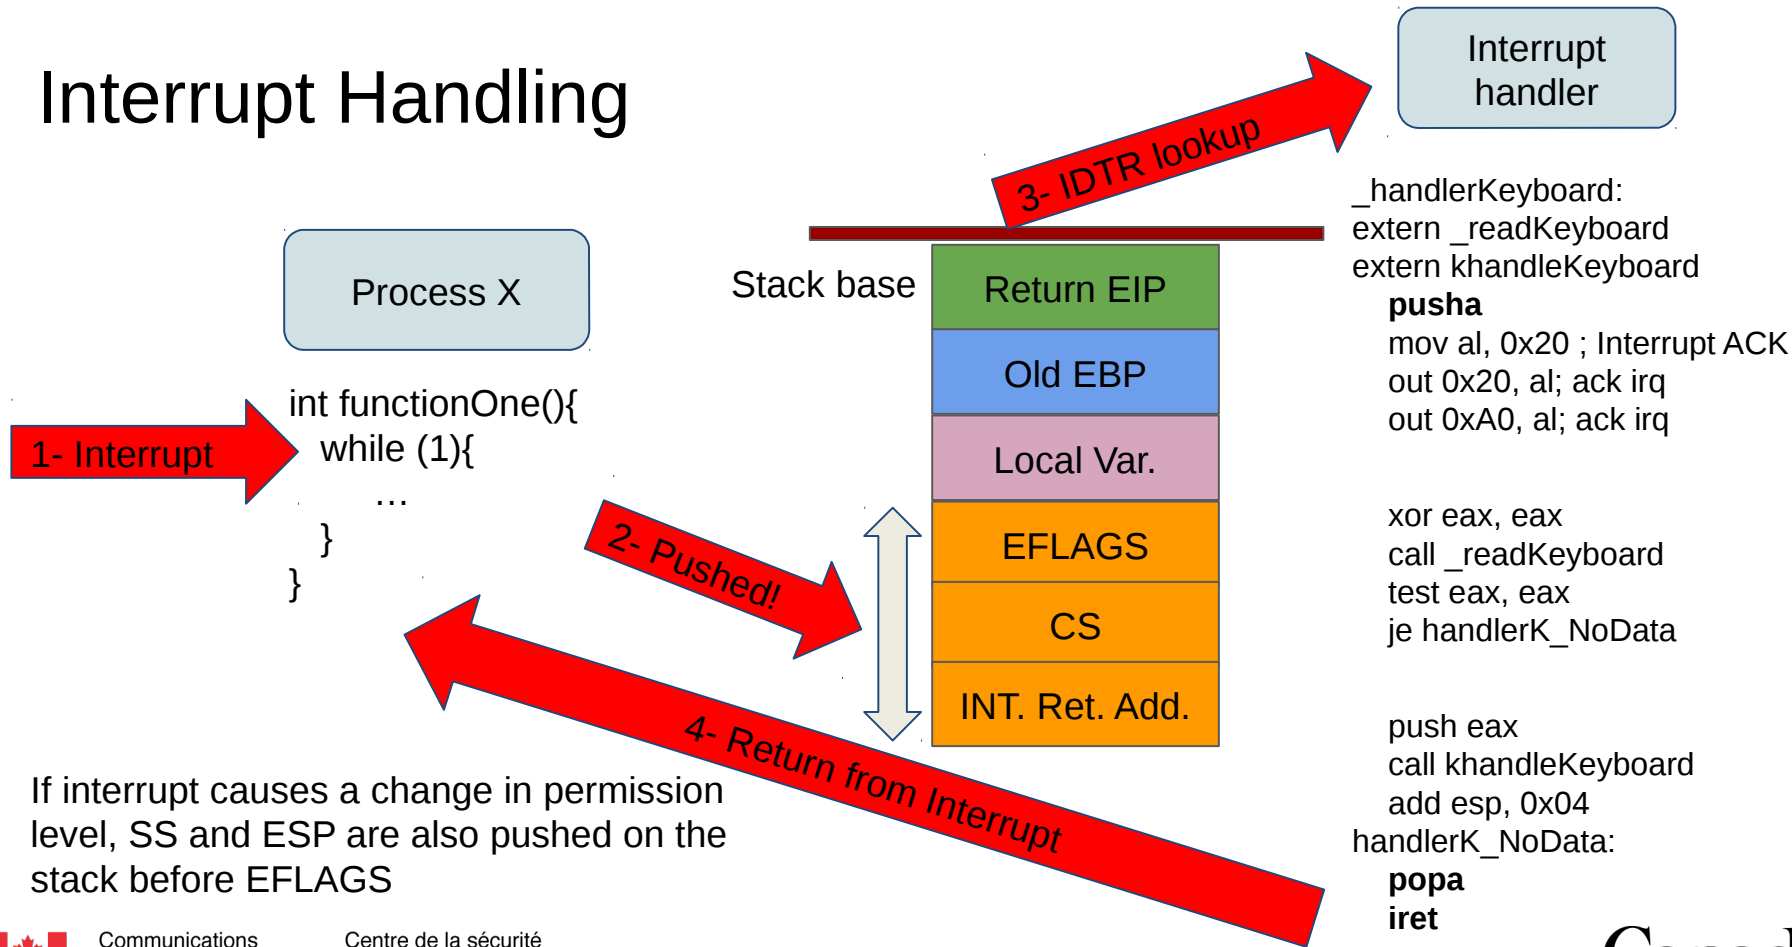


Interrupt delivery to the CPU



Interrupt management stub should be written using assembly.

Interrupt Handling



x86 Exceptions

Exception	Description
0	Divide error
6	Invalid opcode
12	Stack segment fault
14	Page fault
17	Alignment check
19	SIMD floating point exception
...	...



traps.c

```
83 gate_desc idt_table[NR_VECTORS] __page_aligned_bss;
```

On Linux

head_32.s

```
391     movl $idt_table,%edi
392     movl $early_idt_handler_array,%eax
393     movl $NUM_EXCEPTION_VECTORS,%ecx
394 1:
395     movl %eax,(%edi)
396     movl %eax,4(%edi)
397     /* interrupt gate, dpl=0, present */
398     movl $(0x8E000000 + __KERNEL_CS),2(%edi)
399     addl $EARLY_IDT_HANDLER_SIZE,%eax
400     addl $8,%edi
401     loop 1b
402
403     movl $256 - NUM_EXCEPTION_VECTORS,%ecx
404     movl $ignore_int,%edx
405     movl $(__KERNEL_CS << 16),%eax
406     movw %dx,%ax          /* selector = 0x0
407     movw $0x8E00,%dx      /* interrupt gate
408 2:
409     movl %eax,(%edi)
410     movl %edx,4(%edi)
411     addl $8,%edi
412     loop 2b
---
```

segment.h

```
217 #define IDT_ENTRIES          256
218 #define NUM_EXCEPTION_VECTORS 32
```

head_32.s

```
641 idt_descr:
642     .word IDT_ENTRIES*8-1
643     .long idt_table

349     lgdt early_gdt_descr
350     lidt idt_descr
351     ljmp $(__KERNEL_CS),$1f
352 1:     movl $(__KERNEL_DS),%eax
353     movl %eax,%ss
```



irq_vectors.h

```
49 #define IA32_SYSCALL_VECTOR 0x80
```

desc.h

```
506 /*
507  * This routine sets up an interrupt gate at directory privilege level 3
508  */
509 static inline void set_system_intr_gate(unsigned int n, void *addr)
510 {
511     BUG_ON((unsigned)n > 0xFF);
512     _set_gate(n, GATE_INTERRUPT, addr, 0x3, 0, __KERNEL_CS);
513 }
```

traps.c

```
969 void __init trap_init(void)
970 {
971     int i;
972
973 #ifdef CONFIG_EISA
974     void __iomem *p = early_ioremap(0xFFFFD9, 4);
975
976     if (readl(p) == 'E' + ('I' << 8) + ('S' << 16) + ('A' << 24))
977         EISA_bus = 1;
978     early_iounmap(p, 4);
979 #endif
980
981     set_intr_gate(X86_TRAP_DE, divide_error);
982     set_intr_gate_ist(X86_TRAP_NMI, &nmi, NMI_STACK);
983     /* int4 can be called from all */
984     set_system_intr_gate(X86_TRAP_OF, &overflow);
985     set_intr_gate(X86_TRAP_BR, bounds);
986     set_intr_gate(X86_TRAP_UD, invalid_op);
987     set_intr_gate(X86_TRAP_NM, device_not_available);
988
989 #ifdef CONFIG_X86_32
990     set_system_intr_gate(IA32_SYSCALL_VECTOR, entry_INT80_32);
991     set_bit(IA32_SYSCALL_VECTOR, used_vectors);
992 #endif
993 }
```

entry_32.s

```
519 ENTRY(entry_INT80_32)
520     ASM_CLAC
521     pushl    %eax
522     SAVE_ALL pt_regs_ax=-ENOSYS
523
524     /*
525      * User mode is traced as though
526      * turned them off.
527      */
528     TRACE_IRQS_OFF
529
530     movl     %esp, %eax
531     call     do_int80_syscall_32
```



On Linux

common.c

```
335 /* Handles int $0x80 */
336 __visible void do_int80_syscall_32(struct pt_regs *regs)
337 {
338     enter_from_user_mode();
339     local_irq_enable();
340     do_syscall_32_irqs_on(regs);
341 }
```

You are free to continue your exploration from here :)
The point is simply to make you understand a little bit of what happens when a syscall is made on Linux before moving one to understanding how to do a syscall using assembly programming.

