

Let's get modern: 64 bits

Quick tour of “x86-64”

© Government of Canada

This document is the property of the Government of Canada. It shall not be altered, distributed beyond its intended audience, produced, reproduced or published, in whole or in any substantial part thereof, without the express permission of CSE.



Communications
Security Establishment

Centre de la sécurité
des télécommunications

PAGE 1

Canada 

**Reversing 64 bits code is easy:
tracing the code is much easier than
tracing 32 bits code.**



Let's talk registers

- Registers have been extended from 32 to 64 bits
 - EAX -> RAX, ESP -> RSP, etc...
- Some new GPRs are now available
 - R8 to R15
 - 64 bits registers
 - These register can be accessed at byte, word and double word level by appending b, w, d to the register name



Calling convention (functions)

- Function calling convention also changed
- First 6 parameters are now passed in registers
 - Register order goes as: RDI, RSI, RDX, RCX, R8, R9
 - If more than 6 parameters are required, parameters (7th onward) are pushed on the stack in reverse order
- The following registers are “scratch” registers, other registers need to be preserved
 - RAX, RDI, RSI, RDX, RCX, R8, R9, R10, R11
- When calling, RAX is used to indicate the number of vector registers used (can safely ignore for now and simply set RAX to 0 before calling)



16 bytes alignment

- You're also supposed to keep a 16 bytes alignment for the stack
- This means your stack pointer should always end in "0" when looking at it in hexadecimal
- Following the stack frame creation technique should ensure you're 16 bytes aligned:
 - push rbp
 - mov rbp, rsp
- Subtractions to stack pointers should then be on 16 bytes boundaries
 - if you only need 4 bytes, you still sub 16 bytes from the stack pointer
 - Do "sub rsp, 0x10" and NOT "sub rsp, 0x04"
 - This will ensure you're always 16 bytes aligned



But why?



Performance!

- Simply said (omitting a lot of details):
- Cache system will work better with 16 bytes aligned data since a “block” of data will fit the size of a cache line. This will prevent “crossing” cache line boundary in accessing data and will make the performance of memory access much better.
- 16 bytes is also the size of XMM (SIMD) registers.
- On 32 bits systems, this alignment is 4 bytes which makes it almost “forgettable”



Hello 64 bits functions

```
section .data
    hello:
        db "Hello world %d", 0xa, 0

section .text
extern printf

global main

main:

    push rbp
    mov rbp, rsp

    mov rdi, hello
    mov rsi, 0xFF
    xor rax, rax
    call printf

    xor rax, rax
    pop rbp
    ret
```



These rules are true for systems using System V ABI (Linux, MacOS, ...).

Windows uses a different calling convention (RCX, RDX, R8, R9 then stack).

<https://msdn.microsoft.com/fr-ca/library/ms235286.aspx>



Calling convention (system calls)

- System calls calling convention changed
- A **maximum of 6 parameters** is to be used
 - Register order is RDI, RSI, RDX, R10, R8, R9
 - If you forget simply do “man syscall”
- System **call number** is passed into **RAX**
- After syscall, **RAX** contains the **result** of the syscall
- System **calls numbers changed**, full list for 64 bits is located:
 - /usr/include/x86_64-linux-gnu/asm/unistd_64.h
- Syscall now uses the instruction “**syscall**”
 - No more int 0x80



Hello 64 bits system calls

```
section .data
hello:
    db "Hello world",0xa, 0

section .text
extern printf

global main

main:

    push rbp
    mov rbp, rsp

    mov rdi, 0x01 ; stdout
    mov rsi, hello ; buffer
    mov rdx, 0x0c ; buffer size
    mov rax, 0x01 ; syscall number 1 (write)
    syscall

    xor rax, rax
    pop rbp
    ret
```

System V 64 bits ABI can be downloaded here

<https://software.intel.com/sites/default/files/article/402129/mpx-linux64-abi.pdf>

© Government of Canada

This document is the property of the Government of Canada. It shall not be altered, distributed beyond its intended audience, produced, reproduced or published, in whole or in any substantial part thereof, without the express permission of CSE.



Communications
Security Establishment

Centre de la sécurité
des télécommunications

**Obviously other differences exists.
For example, the structures used for
paging are somewhat different. Most
of these differences are of little
interest here.**



Let's code!

